

CS 610 Semester 2024–2025-I

Assignment 2

Divyansh
(210355)

6th September 2024

1 Problem 1

1.1 Performance Bug identification

We start finding performance bug in the `problem.cpp` code by using `perf c2c` hardware counter at a sampling rate of 75000 Hz. Here is the report generated by `perf`:

```
Total Load Opeations      : 6803
Load Local HITM            : 766
Total Shared Cache Lines   : 3
```

Shared Cache Line Distribution Pareto

Cacheline : 0x61734492b380

Local HITM	Data	Address/Offset	Source:Line	NodeCpu List
1.82%		0x8	thread_runner(void*)+817	0{15}
8.07%		0x10	thread_runner(void*)+817	0{0}
20.05%		0x18	thread_runner(void*)+817	0{9}
19.01%		0x20	thread_runner(void*)+817	0{2, 10}
3.91%		0x28	thread_runner(void*)+394	0{0, 2, 9-10, 15}
13.80%		0x30	thread_runner(void*)+830	0{0, 2, 9-10, 12, 15}
15.89%		0x38	pthread_mutex_unlock.c:43	0{0, 2, 9-10, 12, 15}
13.02%		0x38	pthread_mutex_lock.c:48	0{0, 2, 9-10, 12, 15}
2.34%		0x38	lowlevellock.c:45	0{0, 2, 9-10, 12, 15}
2.08%		0x38	lowlevellock.c:42	0{0, 2, 9-10, 12}

Cacheline : 0x61734492b3c0

Local HITM	Data	Address/Offset	Source:Line	NodeCpu List
14.99%		0x0	pthread_mutex_lock.c:94	0{0, 2, 9-10, 12, 15}
17.98%		0x4	pthread_mutex_lock.c:172	0{0, 2, 9-10, 15}
64.03%		0x8	pthread_mutex_lock.c:80	0{0, 2, 9-10, 15}
1.63%		0x8	pthread_mutex_lock.c:44	0{0, 2, 9-10, 12, 15}
1.36%		0x8	pthread_mutex_unlock.c:51	0{0, 2, 9-10, 12, 15}

Table 1: Cache Contention w/o any optimization

1.2 Identifying True Sharing

True Sharing can be seen very evidently in both of the hottest cachelines here. In the first cacheline (0x61734492b380) we can see that for the same cache offsets 0x28, 0x30, 0x38 there are multiple threads contending indicating hits in modified caches. Very similar pattern can be observed in the second cacheline(0x61734492b3c0) at offsets of 0x0, 0x4, 0x8.

CODE One easy observation here which will help in our further analysis that, generally **true sharing** is occurring at the source line of **pthread_mutex_lock** or **unlock** which also makes sense since multiple threads will be contending for the same lock. In code we can find that here:

```

139     while (getline(input_file, line)) {
140
141         pthread_mutex_lock(&line_count_mutex);
142         tracker.total_lines_processed++;
143         pthread_mutex_unlock(&line_count_mutex);
144
145         /* .. some code .. */
146
147         // true sharing: on lock and total word variable
148         pthread_mutex_lock(&tracker.word_count_mutex);
149         tracker.total_words_processed++;
150         pthread_mutex_unlock(&tracker.word_count_mutex);
151
152     }
153 }
```

1.3 Solving True Sharing

This contention can be removed easily. Since we are acquiring lock on every iteration of loop, what we can do instead is that, maintain a local counter for each thread and once the thread is finished reading it will update the counter. Here is how it impacts the performance:

```

Total Load Operations      : 850
Load Local HITM            : 33
Total Shared Cache Lines   : 2
```

Shared Cache Line Distribution Pareto

Cacheline : 0x565585af8380

Local HITM	Data	Address/Offset	Source:Line	NodeCpu	List
68.75%		0x10	thread_runner(void*) + 799	06,	10
31.25%		0x20	thread_runner(void*) + 799	09	

Table 2: Cache Contention w/o true sharing

Analysis We can see that number of modified cache hits has reduced as compared to the unoptimized case. Also, from our observation we can see that the contention for **mutex_lock** didn't come in the radar of **perf** indicating that its number has become too few.

1.4 Identifying False Sharing

In the first cacheline (0x61734492b380) we can see that the core 0{0} writes at the offset of 0x10, the core 0{9} writes at the offset of 0x18 of the same cacheline and the cores 0{2,10} writes at the offset of 0x20 of that exact same cacheline resulting in **false sharing** of this cache line.

CODE Similar to the case of **true sharing** we can observe that **false sharing** mostly occurs at the source line of `thread_runner(void*)+817` where the threads are updating their individual word count in a shared array of `tracker.word_count` which results in contiguous memory of these counters. In code here we can see:

```
148         while ((pos = line.find(delimiter)) != std::string::npos) {
149             token = line.substr(0, pos);
150
151             // false sharing: on word count
152             tracker.word_count[thread_id]++;
153
158
159             line.erase(0, pos + delimiter.length());
160         }
```

Consider the 64 bytes cacheline as an array of 16 blocks each of 4 bytes, then we locate these counters as in the image below which results in conviction of this cacheline from private cache of all 5 threads.

word_counters of all 5 threads

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
---	---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

Figure 1: false sharing of this cacheline because of contiguous memory of counters

1.5 Solving False Sharing

We can remove **False Sharing** by padding memory space in between the counters the threads. Idea is separate the counter memory to different cache lines so that they do not contend for the same cacheline. This can be visualized as in the image below similar to above image. This time the **word_counter** for each thread has been spread across multiple cache lines.

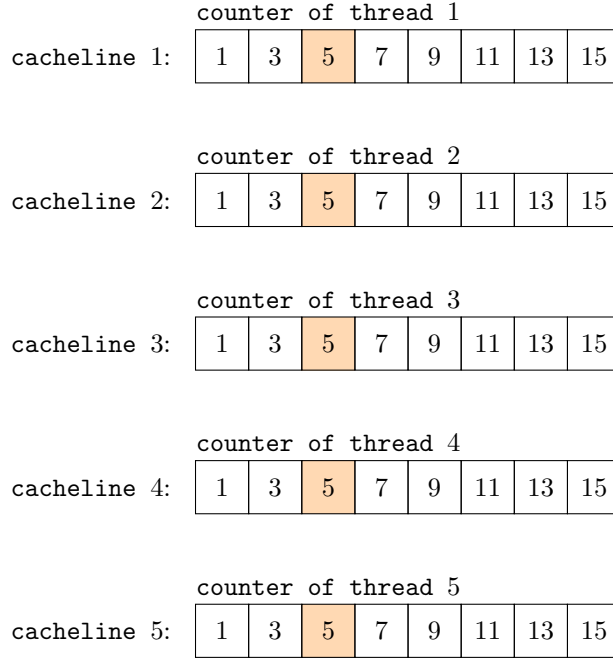


Figure 2: padding to have separate cacheline for each thread counter memory

Employing the above solution, the results from `perf` turns out like following:

```
Total Load Opeations      : 7435
Load Local HITM           : 71
Total Shared Cache Lines   : 2
```

Shared Cache Line Distribution Pareto

Cacheline : 0x5ebc5583e4c0

Local HITM	Data	Address/Offset	Source:Line	NodeCpu List
9.09%		0x0	thread_runner(void*) + 405	0{2, 10, 12, 14}
10.61%		0x10	pthread_mutex_unlock.c:43	0{2, 10, 12, 14}
6.06%		0x10	pthread_mutex_lock.c:48	0{2, 10, 12, 14}
1.52%		0x10	lowlevellock.c:42	0{2, 10, 12, 14}
66.67%		0x20	pthread_mutex_lock.c:80	0{2, 10, 12, 14}
3.03%		0x20	pthread_mutex_unlock.c:51	0{2, 10, 12, 14}
1.52%		0x20	pthread_mutex_lock.c:44	0{2, 10, 12, 14}
1.52%		0x20	pthread_mutex_unlock.c:39	0{2, 10, 12, 14}

Cacheline : 0x5ebc5583e280

Local HITM	Data	Address/Offset	Source:Line	NodeCpu List
100%		0x0	pthread_mutex_unlock.c:80	0{2, 10, 14}

Table 3: Cache Contention w/o false sharing

Analysis We can see that there aren't many different offsets in the cacheline which are being contended by one or two cores. Also, from our observation we can follow that mostly contention are of `mutex.locks` which is `true sharing` issue.

1.6 Combined performance improvement

If remove both `false sharing` and `true sharing` from the source code. Then `perf` report gives empty Pareto table.

```
Total Load Operations      : 1039
Load Local HITM            : 0
Total Shared Cache Lines   : 0
```

```
-----
Shared Cache Line Distribution Pareto
-----
```

Table 4: Cache Contention w/o `false sharing` & `true sharing`

The performance improvement can also be seen the speed ups in completing the task

Case	Time (in ms)	Local HITM
w/o any optimization	399	766
w/o <code>false sharing</code>	326	71
w/o <code>true sharing</code>	181	33
w/o <code>false sharing</code> & <code>true sharing</code>	116	0

Table 5: Speed-up performance

2 Problem 2

Details of this problem is entirely in the source file `problem2.cpp` in the assignment zip folder. Few details of the implementation:

1. The memory buffer is implemented using `queue` data structure with the size limited to `MAX_BUFFER_SIZE`.
2. `pthread_mutex_locks` & `condvars`: One lock on reading the input file to isolate reading. Another lock on writing content to memory buffer. Further, there is a lock and two `condvars` between the consumer and producer threads to maintain correctness.

3 Problem 3

The code we need to analyze is as follows:

```
1  for i = 1, N-2
2      for j = i+1, N
3          A(i, j-i) = A(i, j-i-1) - A(i+1, j-i) + A(i-1, i+j-1)
```

3.1 Dependency Analysis

For all pairs of Array **A** read-write access, we will assume that the write occurred on $(I_o, J_o)^{th}$ iteration and the read happened on $(I_o + \Delta I, J_o + \Delta J)^{th}$ iteration.

1. Dependency in $A(i, j-i)$ & $A(i, j-i-1)$:

$$I_o = I_o + \Delta I \quad \& \quad J_o - I_o = (J_o + \Delta J) - (I_o + \Delta I) - 1$$

This results in $(\Delta I = 0, \Delta J = 1)$ hence the dependency vector is $(0, 1)$. We have the leftmost non-zero value as positive hence it is a **flow dependency**.

2. Dependency in $A(i, j-i)$ & $A(i+1, j-i)$:

$$I_o = I_o + \Delta I + 1 \quad \& \quad J_o - I_o = (J_o + \Delta J) - (I_o + \Delta I)$$

This results in $(\Delta I = -1, \Delta J = -1)$ with distance vector $(-1, -1)$. This is a **anti dependency** both the distances negative.

3. Dependency in $A(i, j-i)$ & $A(i-1, i+j-1)$:

$$I_o = I_o + \Delta I - 1 \quad \& \quad J_o - I_o = (J_o + \Delta J) + (I_o + \Delta I) - 1$$

This results in $(\Delta I = 1, \Delta J = -2I_o)$ hence the dependency vector is $(1, -/ =)$ which means ΔJ can be anything non-positive. Since, the left-most non-zero distance is positive, this is **flow dependency**.