

CS 610 Semester 2024–2025-I: Assignment 2

26th August 2024

Due Your assignment is due by Sep 6, 2024, 11:59 PM IST.

General Policies

- You should do this assignment ALONE.
- Do not copy or turn in solutions from other sources. You will be PENALIZED if caught.

Submission

- Submission will be through Canvas.
- Submit a compressed file called “`<roll>-assign2.tar.gz`”. The compressed file should have the following structure.

```
-- roll
-- -- roll-assign2.pdf
-- -- <problem1-dir>
-- -- -- <source-files>
-- -- <problem2-dir>
-- -- -- <source-files>
```

The PDF file should contain descriptions for the first two problems, and your solution for the last problem.

- We encourage you to use the L^AT_EX typesetting system for generating the PDF file. You can use tools like Tikz, Inkscape, or Draw.io for drawing figures if required. You can alternatively upload a scanned copy of a handwritten solution, but MAKE SURE the submission is legible.
- You will get up to TWO LATE days to submit your assignment, with a 25% penalty for each day.

Evaluation

- Write your programs such that the EXACT output format (if any) is respected.
- We will evaluate your implementations on a Unix-like system, for example, recent Debian-based distributions installed on KD first floor labs.
- We will evaluate the implementations with our OWN inputs and test cases, so remember to test thoroughly.

Problem 1

[40 marks]

You are given a multithreaded program with N threads. The program reads N files, and should report the total number of words and lines processed by all the threads.

We provide a driver source code for the problem. Your task is to analyze the source code, and identify and report the performance bugs present in the source code, if any. If a performance bug is identified, provide a manually fixed version. Describe your modifications to the source code and report the performance gain.

Use the following commands to compile the attached driver and run the sample test case.

```
make
./problem1.out 4 ./test1/input
```

You can use the `perf c2c` tool to identify some forms of performance bugs. You can use the following links to learn more about using `perf`.

- [C2C - False Sharing Detection in Linux Perf](#)
- [perf c2c man page](#)

Input The input to your program will be a path to a file, say `input`.

The file `input` lists the full paths of N source files that are to be processed by N threads. Read the contents of the file `input` into a shared data structure `X`. Each thread will then pick *one* file from the shared data structure to analyze. A thread reads its file one line at a time and divides the line into tokens. The thread updates a counter to track the words encountered by the threads and updates the shared variable to track the total number of places the line.

Example

```
File 1:
ABC, EFG HIJK.
LMNOP QRST.

File 2:
ABC EF HI LMNOPQ
RST UV

Expected Output:
Thread 0 counter: 6
Thread 1 counter: 5
Total words processed: 11
Total lines processed: 4
```

Problem 2

[50 marks]

Write a C++ program that takes five arguments from the command line: a string that represents the path to an input file to be read (say `R`), an integer representing the number of producer threads (say `T`), an integer representing the number of lines each thread should read from the file (say `L`), an integer representing the size of a shared buffer (say `M`), and a string that represents the path to an output file (say `W`).

Assume the file **R** to be read contains **N** lines. The program will launch the required number of threads **T**. AFTER all the threads are created, the threads will **CONTENT** for access to **R** to repeatedly read **L** consecutive lines. Then, each thread will write its share of **L** consecutive lines **ATOMICALLY** to a **FIFO SHARED** buffer. A dedicated consumer thread (not included in **T**) **KEEPS READING** from the shared buffer and writes its contents to the destination file **W**.

- **N** may not be a multiple of **L**. In such cases, the last reader thread will read whatever lines are left.
- Blank lines are also counted.
- The input file can have more than **T*L** lines.
- The problem requires synchronization (barriers, mutexes, and condvars) at multiple places.
- All threads should acquire locks only for the minimum required duration (e.g., reader thread reading **L** lines).
- Use a conditional variable instead of busy waiting for synchronizing accesses to the shared buffer across producer and consumer threads.
- You are not allowed to use concurrent data structures, but you can use other STL data structures like `std::vector` and `std::atomic`.
- The application should terminate properly after all threads are done.

The goal of this problem is to achieve correctness. So, test your code thoroughly with different possible inputs.

Problem 3

[10 marks]

Consider the following loop nest.

```
1  for i = 1, N-2
2    for j = i+1, N
3      A(i, j-i) = A(i, j-i-1) - A(i+1, j-i) + A(i-1, i+j-1)
```

List all flow, anti, and output dependences, if any, using the Delta test. Show your computation. Assume all array subscript references of array **A** are valid.